

Outline for review

Chapter 1 :Introduction

- Internet overview
- what' s a protocol?
- network edge, core, access network
 - packet-switching versus circuit-switching
 - Internet structure
- performance: loss, delay, throughput
- layering, service models
- security
- history

Packet Delay


Transmission delay:

- R = link bandwidth (bps)
- L = packet length (bits)
- Time to send bits into link = L/R

Propagation delay:


- d = length of physical link
- s = propagation speed in medium ($\sim 2 \times 10^8$ m/sec)
- propagation delay = d/s

Note: s and R are very different quantities!



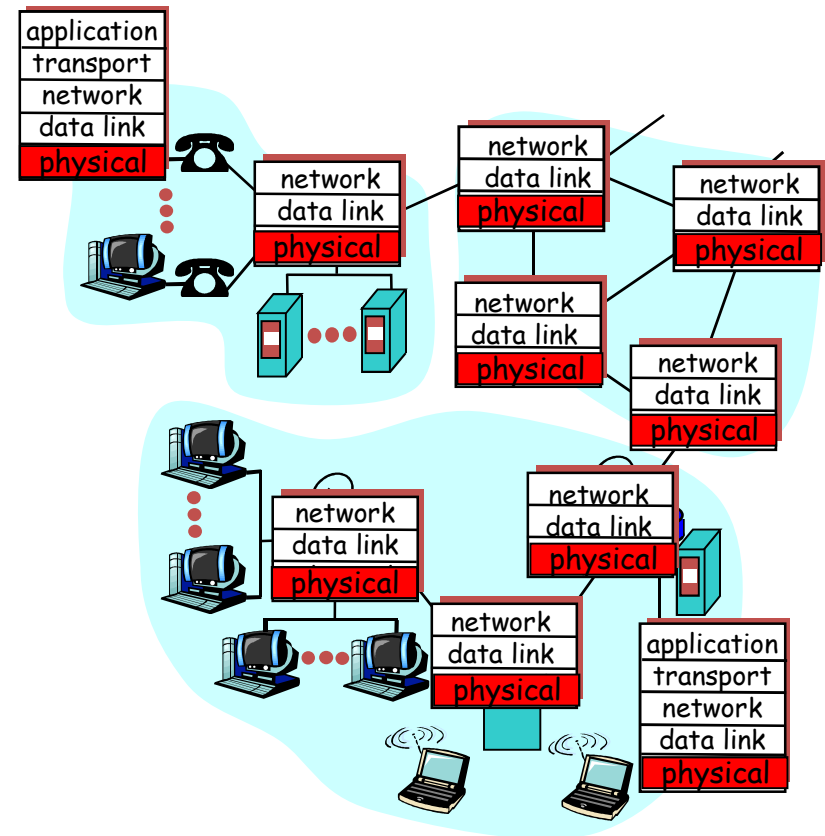
Internet Protocol Stack

- **Application:** supporting network applications
 - FTP, SMTP, HTTP
- **Transport:** process-process data transfer
 - TCP, UDP
- **Network:** routing of datagrams from source to destination
 - IP, routing protocols
- **Link:** data transfer between neighboring network elements
 - PPP, Ethernet
- **Physical:** bits "on the wire"



Chapter 2: Physical Layer

- The function of the physical layer and the issues to be considered
- **Four important characteristics** of the physical layer
- Some Basic Concepts in Data Communication
- Modulation technology and coding technology
- Nyquist's Law and **Shannon's Formula**
- **Multiplexing technology** and **digital signal coding technology**
- Commonly used transmission medias



Channel Characteristics

- Shannon's Formula: limit information transmission rate C of channel can be expressed as:

$$C = W \log_2(1 + S/N)$$

- W - Frequency bandwidth in Hz
- S - Average signal power through the channel
- N - Gaussian noise power through the channel
- S/N – related to signal-to-noise ratio
 - Typically in db (decibels): $10 \log_{10} S/N$
- Actual rate much lower than C due to signal loss
- Possible to achieve error-free transmission as long as transmission rate $< C$

Digital signal coding technology

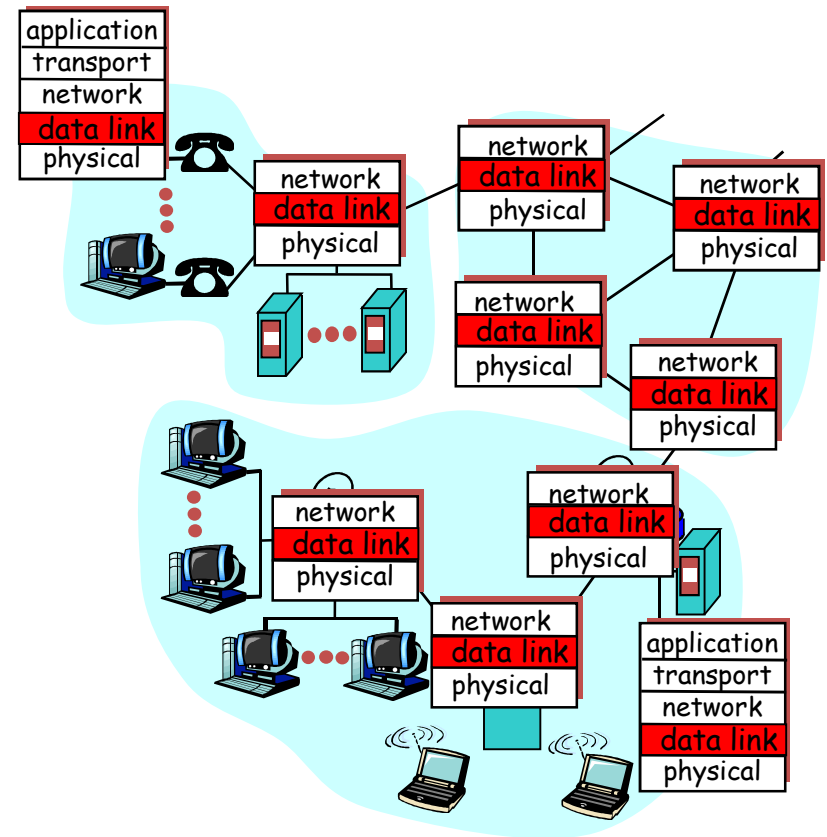
- Non-Return-to-Zero (NRZ)
- Manchester encoding
- Differential Manchester encoding

Exercise:

Draw the waveform of "001101" with NRZ, Manchester encoding, differential Manchester encoding.

Chapter 3: the data link layer

- principles behind data link layer services:
 - framing
 - error detection, correction
 - reliable data transfer
 - sharing a broadcast channel: media access control (MAC)

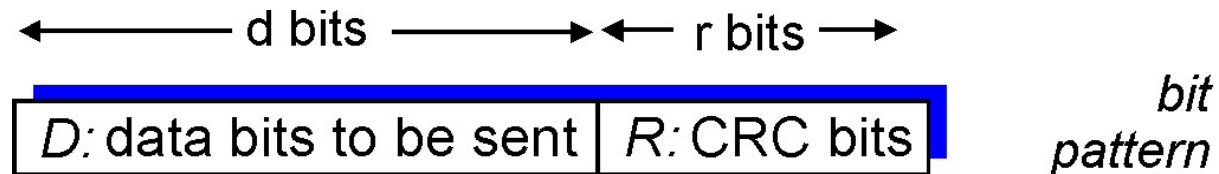


Framing

- **Character count method :**
 - encapsulate datagram into frame, adding header (character number)
- **First and tail bound method based on character:**
 - encapsulate datagram into frame, adding header, trailer (character)
- **First and tail bound method based on bit:**
 - encapsulate datagram into frame, adding header, trailer (bit sequence, e.g. 7EH)
- **Physical layer coding violation method:**
 - encapsulate datagram into frame **without stuffing**
 - Only be used in the networks with redundancy coding technology in the physical layer

Checksumming: Cyclic Redundancy Check

- view data bits, **D**, as a binary number
- choose $r+1$ bit pattern (generator), **G**
- goal: choose r CRC bits, **R**, such that
 - $\langle D, R \rangle$ exactly divisible by G (modulo 2)
 - receiver knows G , divides $\langle D, R \rangle$ by G . If non-zero remainder: error detected!
 - can detect all burst errors less than $r+1$ bits
- widely used in practice (ATM, HDCL)



$$D * 2^r \text{ XOR } R$$

mathematical formula

Principles of reliable data transfer

- From rdt 1.0 to rdt 3.0: **Assumptions and mechanisms**
 - rdt 1.0: over a perfectly reliable channel
 - rdt2.x (rdt2.0, 2.1, 2.2): channel with bit errors
 - rdt3.0: channels with errors and loss
- Pipelined protocols
 - Go-back-N
 - Selective Repeat

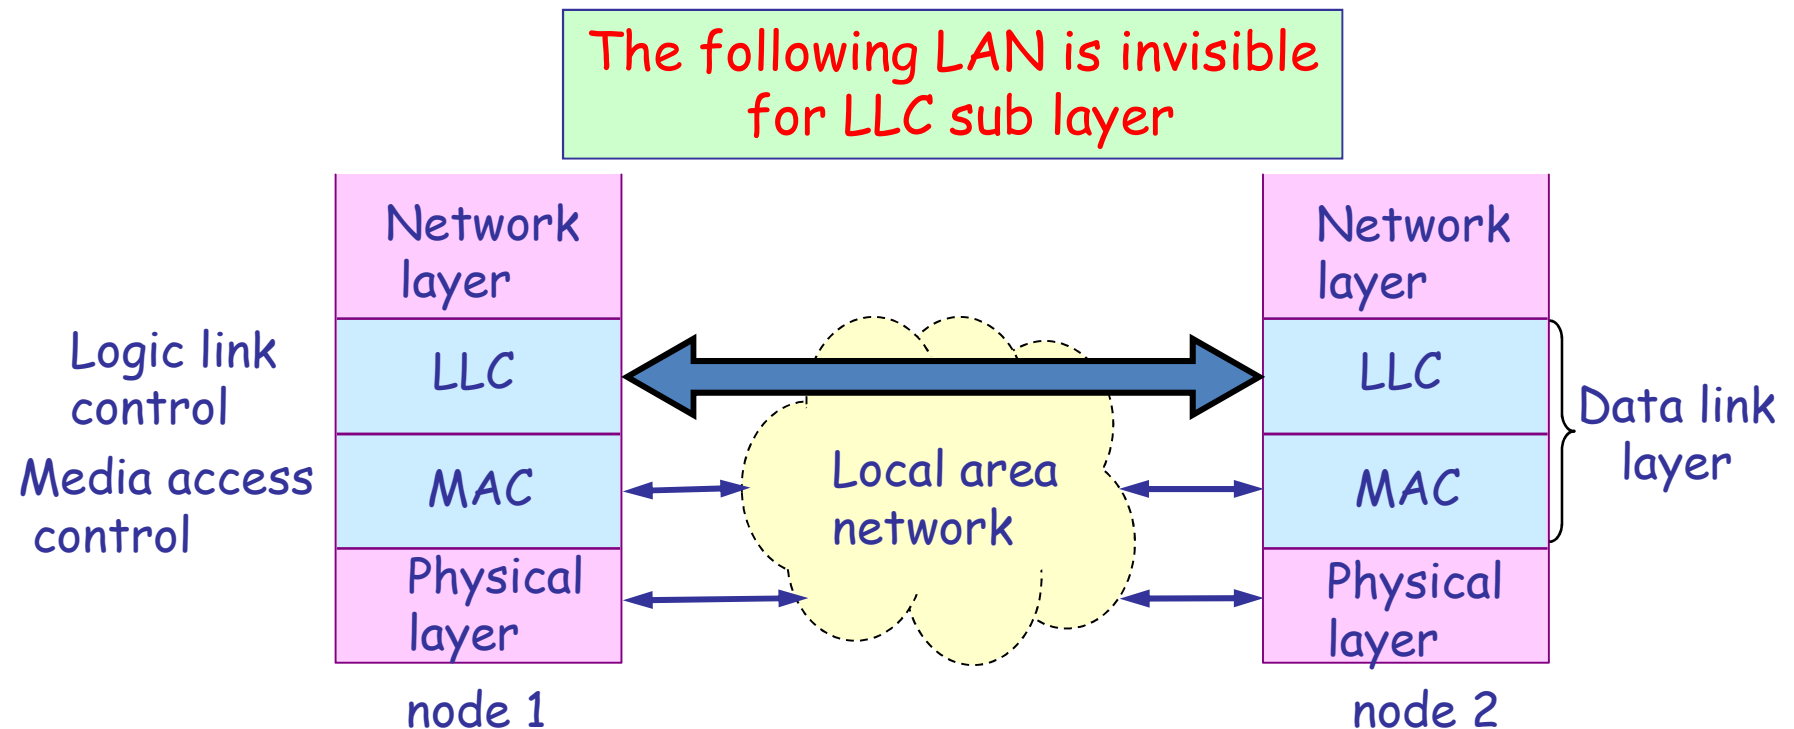
Summary of MAC protocols

- What do you do with a shared media?
 - **Channel Partitioning**, by time, frequency or code
 - Time Division, Code Division, Frequency Division
 - **Random partitioning** (dynamic),
 - ALOHA, S-ALOHA, CSMA, CSMA/CD
 - carrier sensing: easy in some technologies (wire), hard in others (wireless)
 - CSMA/CD used in Ethernet
 - **Taking Turns**
 - polling from a central cite, token passing

Chapter 4: Local Area Network

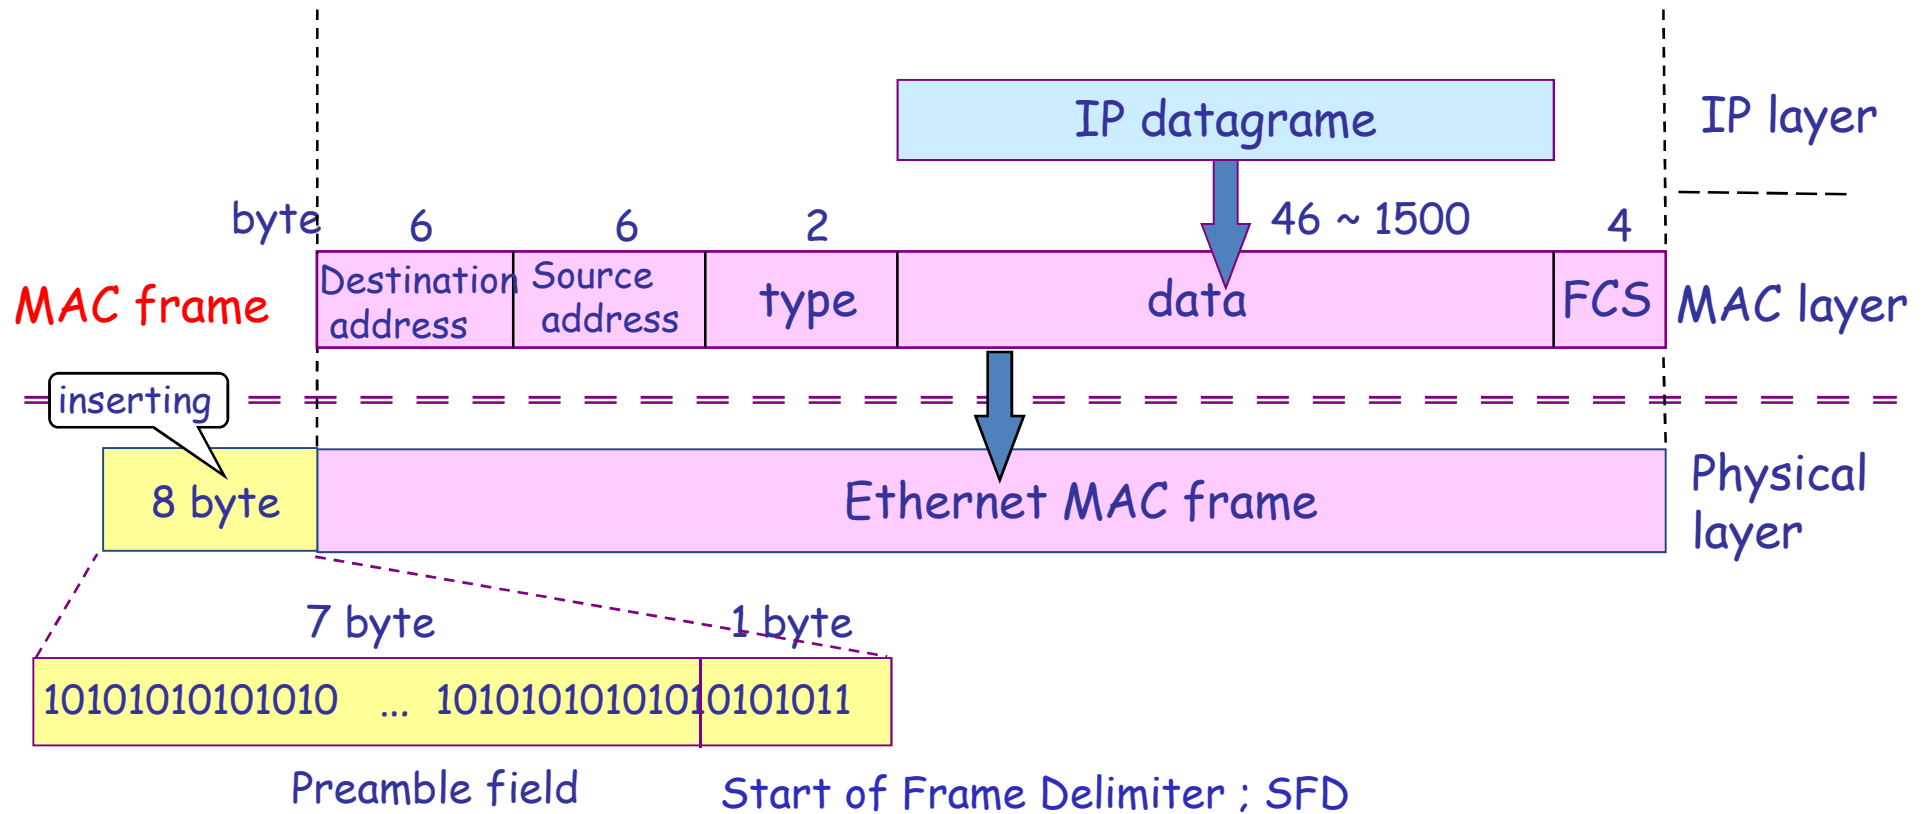
- LAN model
- Ethernet /IEEE 802.3(Ethernet V2 frame structure, CSMA/CD, MAC address)
- hubs, bridges, switches
- WLAN (IEEE 802.11, CSMA/CA, hidden terminal problem, RTS/CTS)

LAN model



For the same LLC, several MAC options may be provided.

Ethernet Frame Structure



Ethernet: uses CSMA/CD

A: sense channel, if idle

then {

transmit and monitor the channel;

If detect another transmission

then {

abort and send jam signal;

update # collisions;

delay as required by exponential backoff algorithm;

goto A

}

else {done with the frame; set collisions to zero}

}

else {wait until ongoing transmission is over and goto A}

Ethernet's CSMA/CD (more)

Exponential Backoff:

- *Goal*: adapt retransmission attempts to estimated current load
 - heavy load: random wait will be longer
- first collision: choose K from {0,1}; delay is K x 512 bit transmission times
- after second collision: choose K from {0,1,2,3}...
- after ten or more collisions, choose K from {0,1,2,3,4,...,1023}

Hubs, switches and routers


- Hubs: **Physical Layer** devices, **do not isolate** collision domains, (results in no increase in max throughput), cannot connect different Ethernet types (e.g., 10BaseT and 100baseT)
- Bridges: **Link Layer** devices, **isolates collision domains** since it buffers frames; bridges **filter** packets, **do not isolate broadcast domains**; maintain filtering tables, implement filtering, learning and spanning tree algorithms
- routers: **network layer** devices (examine network layer headers), **isolates broadcast domains** ; maintain routing tables, implement routing algorithms

802.11 MAC

- MAC layer covers three functional areas:
 - Reliable data delivery
 - ACK-based scheme for reliability (receiver sends ACK after each successful transmission)
 - Medium access control
 - **CSMA/CA**; collision avoidance, not collision detection, **Why? How?**
 - Security
 - Wired Equivalent Privacy (WEP), WEP relies on a secret key being shared by end hosts and APs


Wireless network characteristics

Multiple wireless senders and receivers create additional problems (beyond multiple access):



Hidden terminal problem

- ❖ B, A hear each other
- ❖ B, C hear each other
- ❖ A, C can not hear each other means A, C unaware of their interference at B



Signal attenuation:

- ❖ B, A hear each other
- ❖ B, C hear each other
- ❖ A, C can not hear each other interfering at B

Avoiding collisions (more)

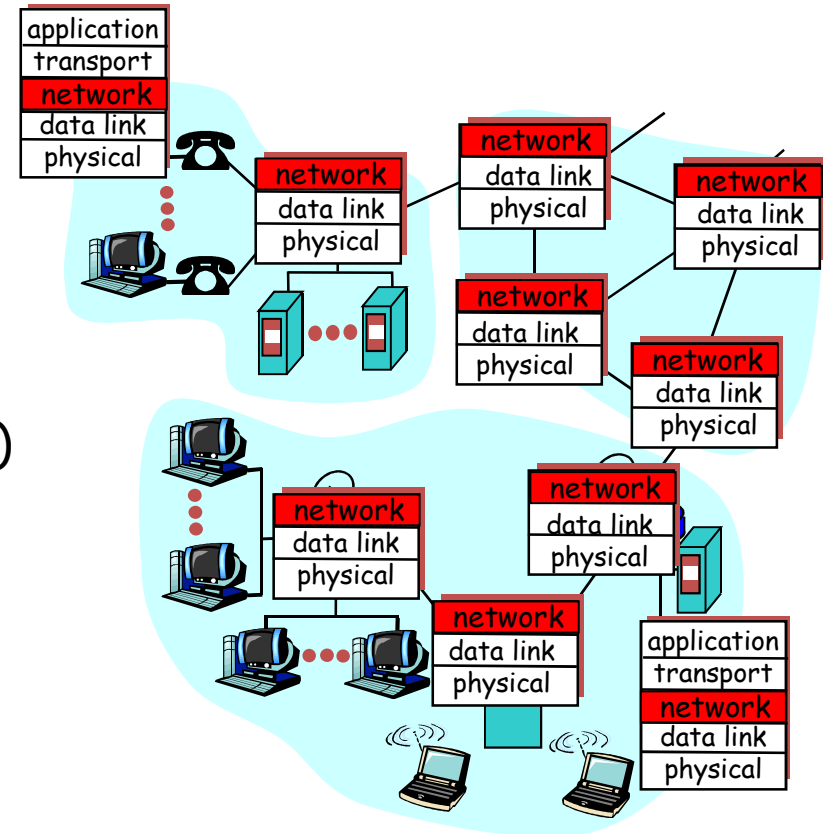
idea: allow sender to “reserve” channel rather than random access of data frames: avoid collisions of long data frames

- sender first transmits *small* request-to-send (RTS) packets to BS using CSMA
 - RTSs may still collide with each other (but they’re short)
- BS broadcasts clear-to-send CTS in response to RTS
- CTS heard by all nodes
 - sender transmits data frame
 - other stations defer transmissions

*avoid data frame collisions completely
using small reservation packets!*

Chapter 5: Network Layer

- Virtual circuit and datagram networks
- Routing algorithms:
 - Dijkstra's algorithm
 - Broadcast routing
 - Link state
 - Distance vector ("count to infinity" problem)
- Routing in the Internet (RIP, OSPF, BGP)
- IP: Internet protocol
 - IPv4 Datagram format
 - IP fragment
 - IPv4 addressing
 - NAT
 - ARP
 - ICMP
 - IPv6
 - From IPv4 to IPv6



Link state algorithm

Broadcast routing + Dijkstra's algorithm

- Having each node **broadcast** link-state packets to **all** other nodes → all nodes have an identical and complete view of the network.
- Using Dijkstra's algorithm **compute** the least-cost path from one node to all other nodes in the network
- If a link cost changes, re-broadcast and re-compute

Dijkstra's algorithm

1 **Initialization:**

2 $N' = \{u\}$

3 for all nodes v

4 if v adjacent to u

5 then $D(v) = c(u,v)$

6 else $D(v) = \infty$

7

8 **Loop**

9 find w not in N' such that $D(w)$ is a minimum

10 add w to N'

11 update $D(v)$ for all v adjacent to w and not in N' :

12 $D(v) = \min(D(v), D(w) + c(w,v))$

13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in N'**



Distance vector algorithm

Iterative, asynchronous:

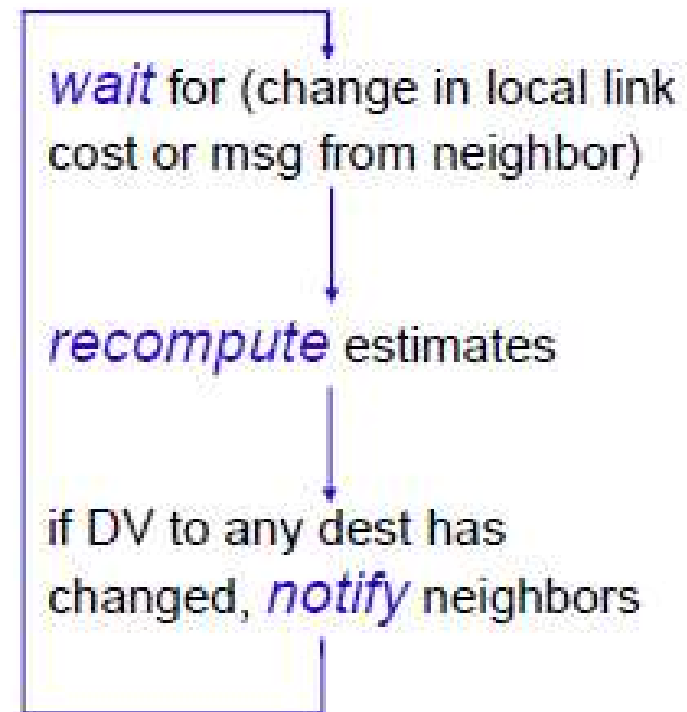
Each local iteration caused by:

- Local link cost change
- DV update message from neighbor

Distributed:

- Each node notifies neighbors **only when** its DV change
- Neighbors then notify their neighbors if necessary

Each node:



Count to infinity problem

Link cost changes:

- good news travels fast
- **Bad news** travels slow -
"count to infinity"
problem!

How to solve the count to infinity problem:

- **Poisoned reverse**

"bad
news
travels
slow"

"count to
infinity"
problem



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

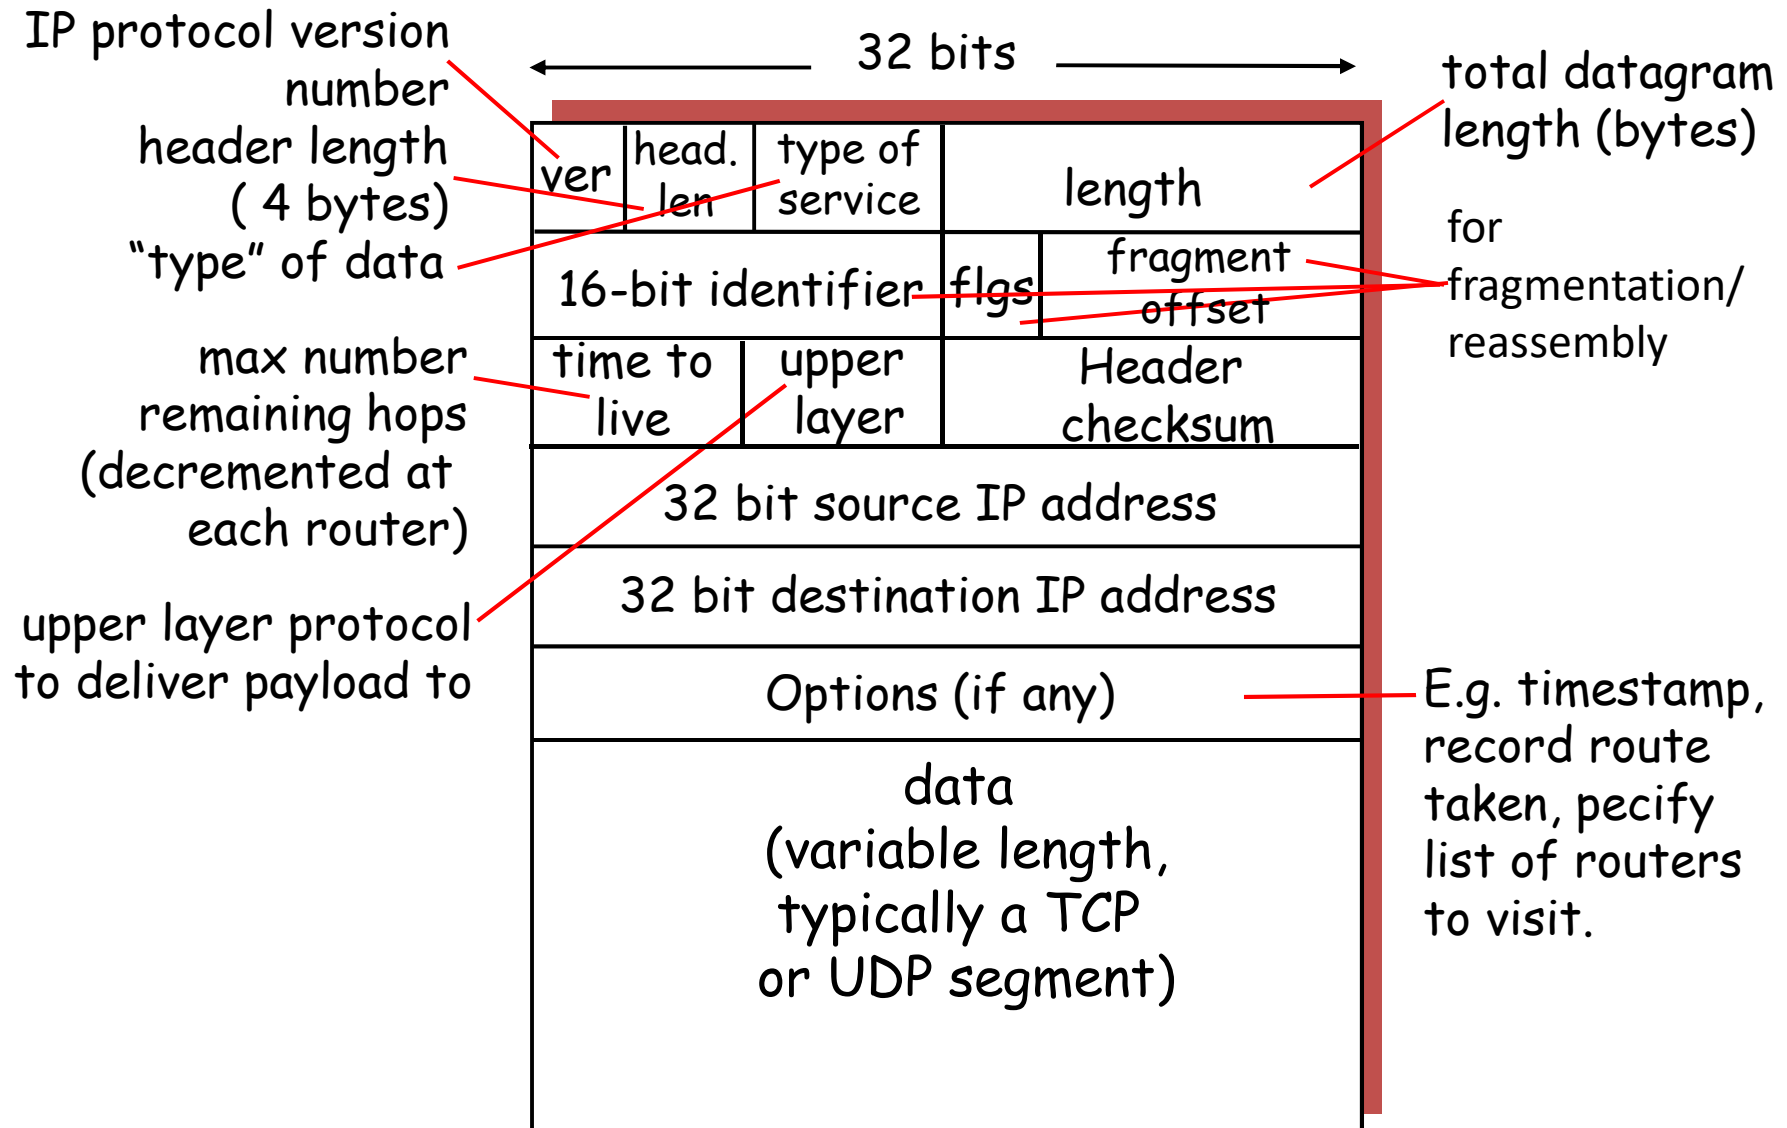
node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

node z table


		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

IP datagram format



IP Fragmentation & Reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame.
 - different link types, different MTUs
- large IP datagram divided ("fragmented") within net
 - one datagram becomes several datagrams
 - "reassembled" only at final destination
 - IP header bits used to identify, order related fragments



Q: A datagram of 4000 bytes (20 bytes of IP header) arrives at a router and must be forwarded to a link with an MTU of 1500 bytes. How to do?

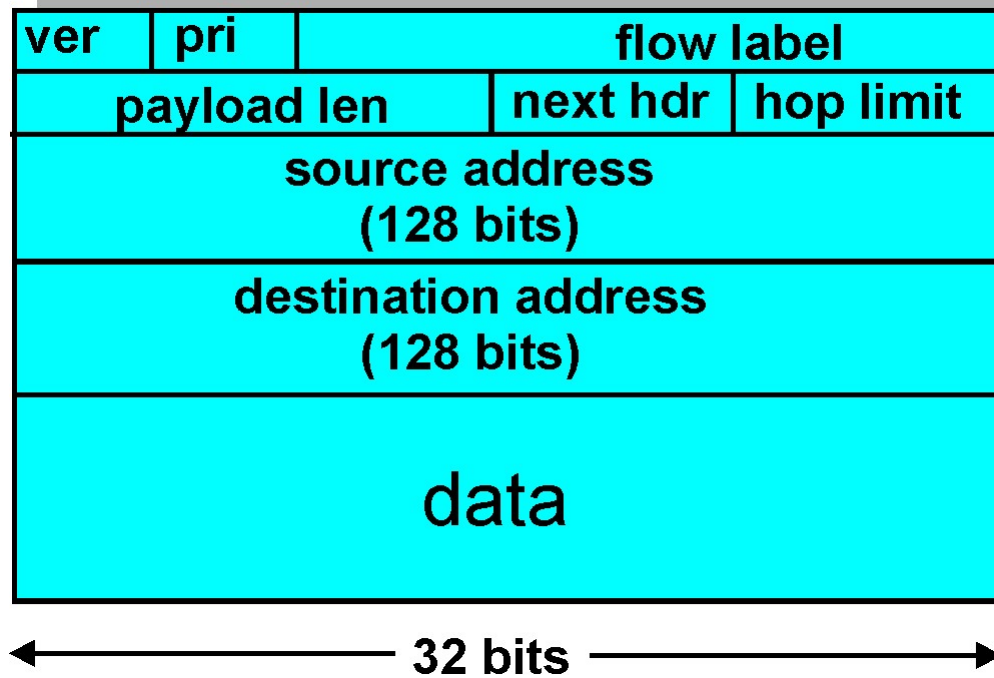
IPv4 addressing

- Dotted-decimal notation: 193.32.216.9
- classful addressing:
- Subnetting and subnet mask
- **CIDR: Classless InterDomain Routing**
 - network portion of address of arbitrary length
 - address format: **a.b.c.d/x**



200.23.16.0/23

IPv6 Header



IPv6 address

- Three types: unicast, multicast, **anycast**
- **Colon hexadecimal notation:**

68E6:8C64:FFFF:FFFF:0:1180:960A:FFFF

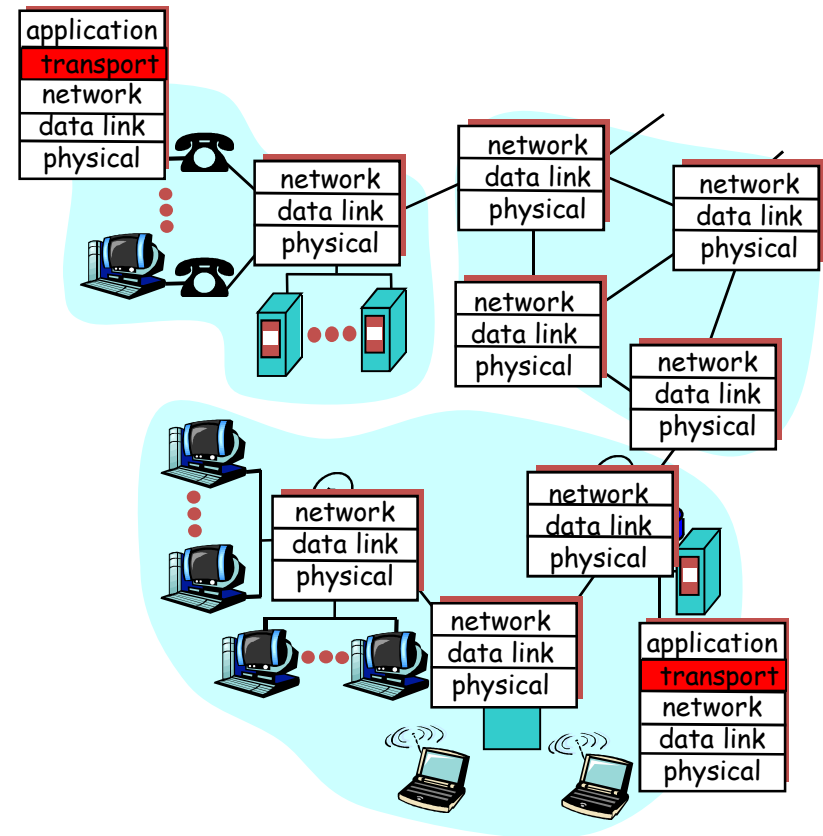
- **Zero compression:**

FF05:0:0:0:0:0:0:0:B3 == FF05::B3 ;

0:0:0:0:0:0:128.10.2.1 == ::128.10.2.1

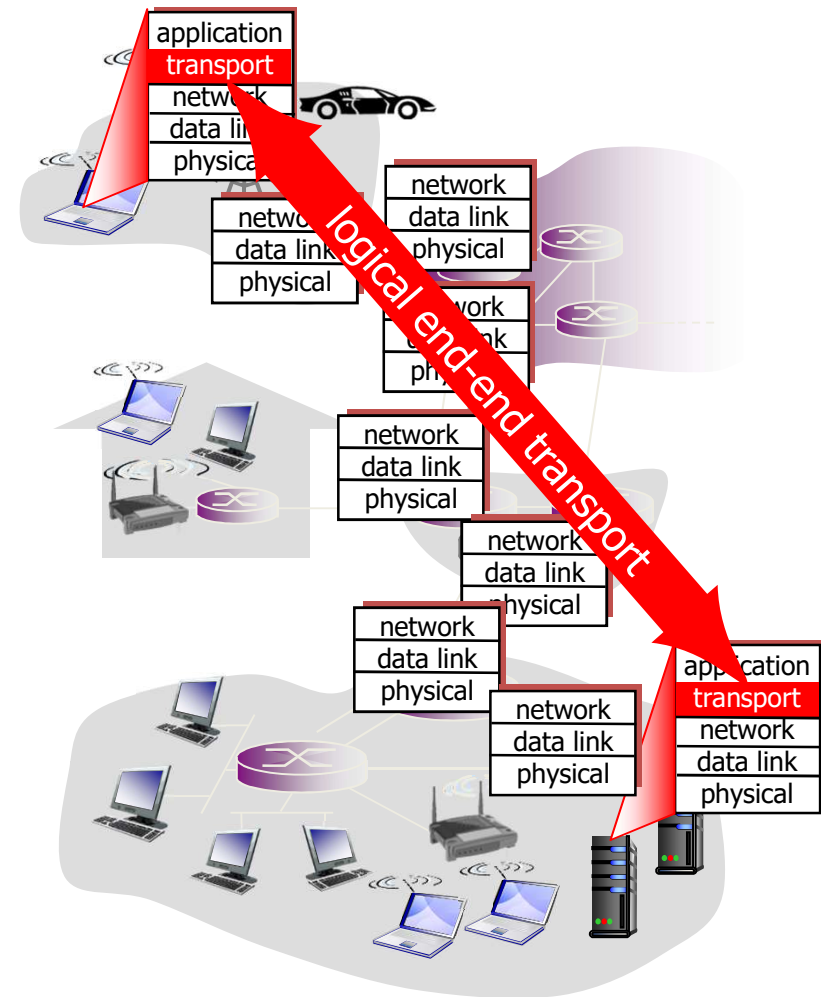
Chapter 6: Transport Layer

- Principles behind transport layer services:
 - Multiplexing, demultiplexing: port numbers and sockets
 - Reliable data transfer
 - Flow control
 - Congestion control
- Two transport layer protocols in the Internet
 - TCP, UDP
 - How to choose?



Internet transport-layer protocols

- reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: UDP
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



TCP: Overview

- point-to-point:
 - one sender, one receiver
- **reliable, in-order byte stream:**
 - no “message boundaries”
- **pipelined:**
 - TCP congestion and flow control set window size
- **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- **connection-oriented:**
 - handshaking (exchange of control msgs) initializes sender, receiver state before data exchange
- **flow controlled:**
 - sender will not overwhelm receiver
- **Congestion controlled**

TCP seq. numbers, ACKs

sequence numbers:

- byte stream “number” of first byte in segment’s data

acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

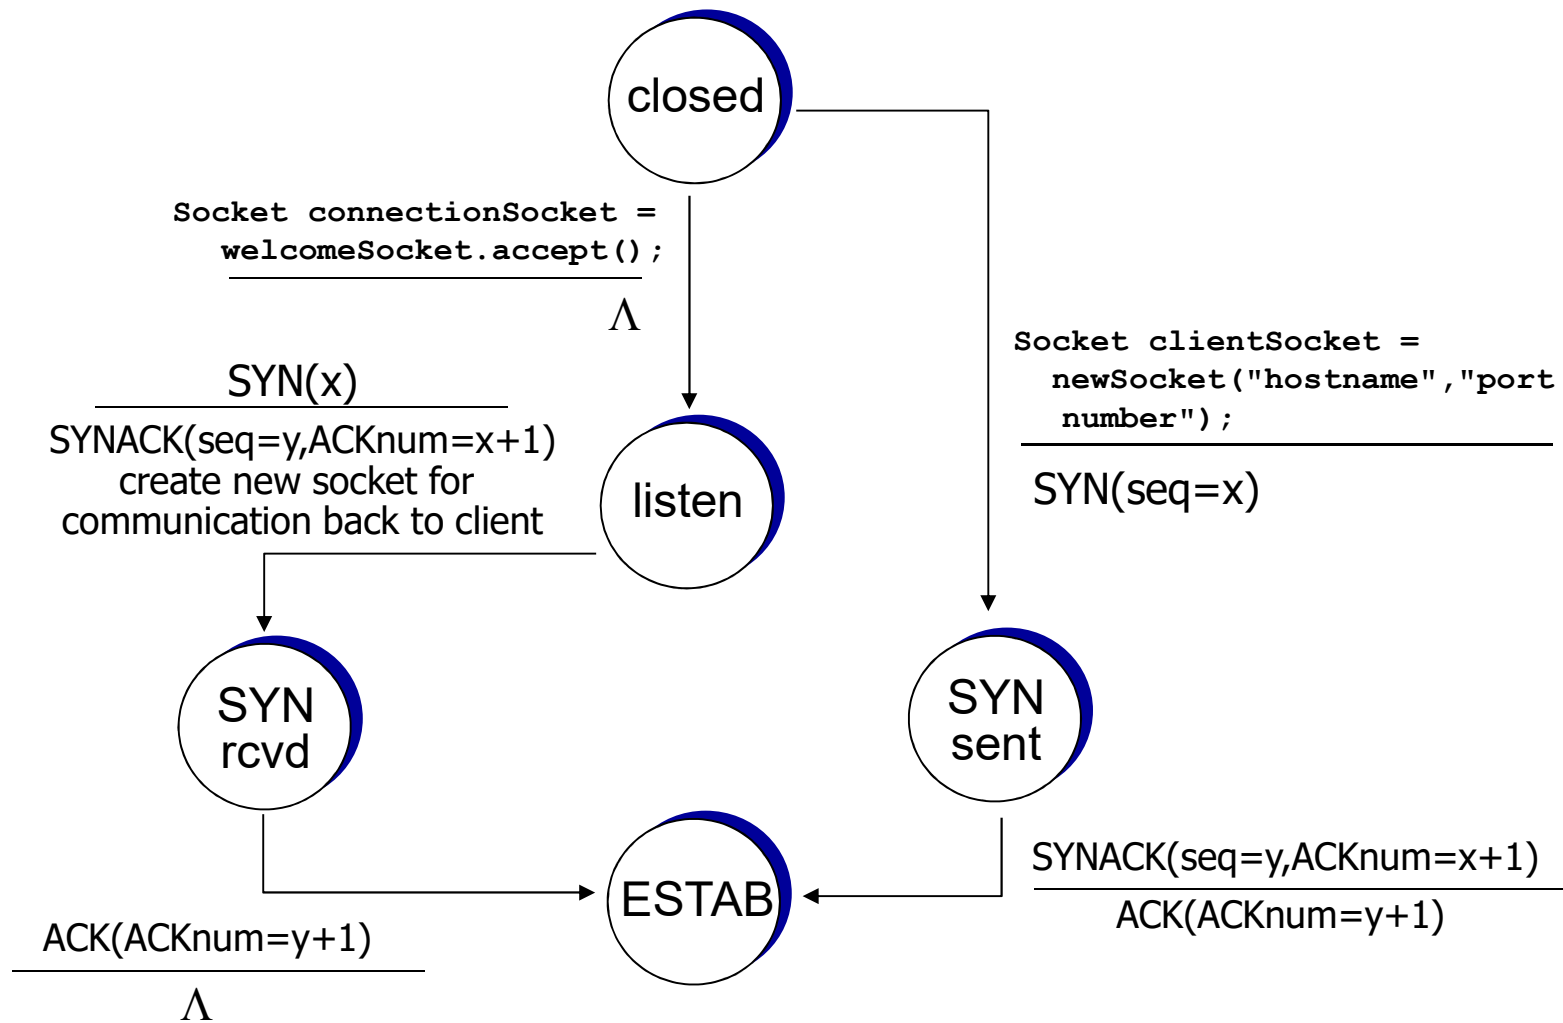
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

TCP 3-way handshake: FSM




TCP congestion control:

additive increase multiplicative decrease


- *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase **cwnd** by 1 MSS every RTT until loss detected
 - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth



TCP Slow Start

- when connection begins, increase rate exponentially until first loss event:
 - initially **cwnd** = 1 MSS
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- summary: initial rate is slow but ramps up exponentially fast



TCP: detecting, reacting to loss

- loss indicated by timeout:
 - **cwnd** set to 1 MSS;
 - window then grows exponentially (as in slow start) to threshold, then grows linearly
- loss indicated by 3 duplicate ACKs: TCP RENO
 - dup ACKs indicate network capable of delivering some segments
 - **cwnd** is cut in half window then grows linearly
- TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate acks)

